

ELENA ROSTOVA

San Francisco, CA | elena.rostova@example.com | github.com/erostova

PROFESSIONAL SUMMARY

Senior AI Systems Engineer with 6.5 years of experience architecting on-device inference engines and distributed model serving pipelines. Specializing in LiteRT, vLLM, and low-latency quantization.

EXPERIENCE

Staff Edge AI Engineer | NeuroEdge Labs (2022 - Present)

- Optimized on-device LLM inference using Google LiteRT and INT4 weight quantization, achieving a 45% latency reduction on Qualcomm Hexagon NPUs.
- Orchestrated distributed model serving clusters across 120 Kubernetes nodes handling over 50,000 QPS with 99.99% uptime.
- Deployed vLLM with PagedAttention and continuous batching, cutting cloud GPU infrastructure costs by \$1.2M annually.

ELENA ROSTOVA (Page 2)

Senior Systems Engineer | ScaleStream (2019 - 2022)

- Built high-throughput gRPC microservices in C++ and Python for real-time video analytics.
- Automated CI/CD deployment pipelines using Docker, Kubernetes (k8s), and Helm.

EDUCATION

Bachelor of Science in Computer Science | UC Berkeley (2015 - 2019)

TECHNICAL SKILLS

Languages: Python, C++, Rust

Technologies: LiteRT, vLLM, Kubernetes, Docker, Triton, ONNX, GGUF, AWS